



# Clean no-tool comparison on eight exposed benchmark tasks

**A Normalized task score**

|                   | Knowledge | Drug retrieval | Target retrieval | Property prediction | Virtual screening | Binding mechanism | Molecular design | Retro-synthesis |
|-------------------|-----------|----------------|------------------|---------------------|-------------------|-------------------|------------------|-----------------|
| GPT-5.4           | 0.10      | 0.35           | 0.30             | 0.54                | 0.32              | 0.19              | 0.90             | 0.05            |
| GPT-5.5           | 0.15      | 0.48           | 0.28             | 0.62                | 0.37              | 0.28              | 0.94             | 0.47            |
| GPT-5.6 Sol       | 0.15      | 0.15           | 0.32             | 0.77                | 0.42              | 0.33              | 0.93             | 0.50            |
| DeepSeek V4 Flash | 0.15      | 0.16           | 0.28             | 0.56                | 0.00              | 0.24              | 0.74             | 0.50            |
| DeepSeek V4 Pro   | 0.05      | 0.64           | 0.32             | 0.41                | 0.11              | 0.27              | 0.93             | 0.00            |
| Kimi K2.5         | 0.15      | 0.30           | 0.28             | 0.75                | 0.05              | 0.24              | 0.95             | 0.09            |
| Qwen3.7 Plus      | 0.05      | 0.00           | 0.25             | 0.69                | 0.11              | 0.30              | 0.95             | 0.25            |
| GLM 5.2           | 0.10      | 0.05           | 0.23             | 0.55                | 0.26              | 0.07              | 0.91             | 0.47            |
| MiniMax M3        | 0.45      | 0.17           | 0.28             | 0.74                | 0.16              | 0.40              | 0.94             | 0.10            |
| MiMo V2.5         | 0.05      | 0.24           | 0.34             | 0.55                | 0.11              | 0.17              | 0.73             | 0.42            |



**B Macro mean  
95% task bootstrap CI**



Each cell aggregates 20 author-supplied cases. Gold answers were withheld during inference. GPT models used low reasoning; five OpenRouter reasoning-first models used the preregistered no-reasoning compatibility amendment.